

CMSC 722, Spring 2018, Project 1 (Revised):


Compare the FF heuristic with a domain-specific heuristic

Last update February 28, 2018

- ▶ To be done individually (not in teams)
- Due date: March 14
- Late date (10% off): March 16

Problem domain

- Modified version of [Racetrack](#)
 - Invented in early 1970s
 - played by hand on graph paper
- 2-D polygonal region
 - Inside are a starting point, finish line, maybe obstacles
- All walls are straight lines
- All coordinates are nonnegative integers
- Robot vehicle begins at starting point, can make certain kinds of moves
- Want to move it to the finish line as quickly as possible
 - Without crashing into any walls
 - Need to come to a complete stop on the finish line




A domain-specific representation

- Later I'll discuss some state-variable representations
- Current state $s_{i-1} = (p_{i-1}, z_{i-1})$
 - ▶ location $p_{i-1} = (x_{i-1}, y_{i-1})$, nonnegative integers
 - ▶ velocity $z_{i-1} = (u_{i-1}, v_{i-1})$, integers
- To move the vehicle
 - ▶ First choose a new velocity $z_i = (u_i, v_i)$, where

$$u_i \in \{u_{i-1} - 1, u_{i-1}, u_{i-1} + 1\}, \quad (1)$$

$$v_i \in \{v_{i-1} - 1, v_{i-1}, v_{i-1} + 1\}. \quad (2)$$

- ▶ New location: $p_i = (x_{i-1} + u_i, y_{i-1} + v_i)$
- ▶ New state: $s_i = (p_i, z_i)$



Example

- Initial state:

$$p_0 = (4, 5)$$

$$z_0 = (0, 0)$$

$$s_0 = (p_0, z_0) = ((4, 5), (0, 0))$$

- First move:

$$z_1 = (0, 0) + (1, 1) = (1, 1)$$

$$p_1 = (4, 5) + (1, 1) = (5, 6)$$

$$s_1 = ((5, 6), (1, 1))$$

- Second move:

$$z_2 = (1, 1) + (-1, 1) = (0, 2)$$

$$p_2 = (5, 6) + (0, 2) = (5, 8)$$


$$s_2 = ((5, 8), (0, 2))$$

- Third move:

$$z_3 = (0, 2) + (1, 0) = (1, 2)$$


$$p_3 = (5, 8) + (1, 2) = (6, 10)$$

$$s_3 = ((6, 10), (1, 2))$$



Walls

- *edge*: a pair of points (p, q)
 - $p = (x, y), q = (x', y')$
 - coordinates are nonnegative integers
- *wall*: an edge that the vehicle can't cross



- List of walls in the example:

$[((0, 0), (10, 0)), ((10, 0), (10, 10)), ((10, 10), (20, 10)),$
 $((20, 10), (30, 0)), ((30, 0), (30, 10)), ((30, 10), (10, 20)),$
 $((10, 20), (0, 20)), ((0, 20), (0, 0)), ((3, 14), (10, 14)),$
 $((10, 14), (10, 16)), ((10, 16), (3, 16)), ((3, 16), (3, 14))]$

Moves and paths

- *move*: an edge $m = (p_{i-1}, p_i)$
 - ▶ $p_{i-1} = (x_{i-1}, y_{i-1})$
 - ▶ $p_i = (x_i, y_i)$
 - ▶ represents change in location from time $i - 1$ to time i


- Example:

$$m_1 = ((4, 5), (5, 6))$$

$$m_2 = ((5, 6), (5, 8))$$

$$m_3 = ((5, 8), (6, 10))$$


$$m_4 = ((6, 10), (8, 12))$$



- *path*: list of locations $[p_0, p_1, p_2, \dots, p_n]$
 - ▶ represents sequence of moves $(p_0, p_1), (p_1, p_2), (p_2, p_3), \dots, (p_{n-1}, p_n)$
 - ▶ Example: $[(4, 5), (5, 6), (5, 8), (6, 10), (8, 12)]$
- If a move or path intersects a wall, it *crashes*, otherwise it is *safe*

Objective

- *Finish line*:
 - ▶ an edge $f = ((q, r), (q', r'))$
 - ▶ always horizontal or vertical
- Want to reach the finish line as quickly as possible
 - ▶ as few moves as possible
- Find a path $[p_0, p_1, \dots, p_n]$
 - ▶ p_n must be on the finish line
 - $\exists t, 0 \leq t \leq 1$, such that $p_n = t(q, r) + (1 - t)(q', r')$
 - ▶ Final velocity must be $(0, 0)$
 - Thus $p_{n-1} = p_n$



Example: $[(4, 5), (5, 6), (5, 8), (6, 10), (8, 12), (11, 13), (15, 13), (19, 12), (22, 11), (24, 10), (25, 9), (25, 8), (25, 8)]$

Three state-variable representations

- Unlike the state-variable representations in the book,
 - value of a state variable can be an arbitrary data structure
 - preconditions, effects, goal tests, rigid conditions can be computational formulas
 - ▶ OK since we're doing forward search
- Representation 1: state variables \mathbf{x} , \mathbf{y} for location, and \mathbf{u} , \mathbf{v} for velocity
 - ▶ Each state variable's value is an integer
- Representation 2: state variables \mathbf{p} and \mathbf{z} : current location, current velocity
 - ▶ Each state variable's value is a pair of integers
- Representation 3: one state variable \mathbf{s} : current state
 - ▶ Value is a 4-tuple of integers
- In each representation, the locations of the walls are rigid properties
 - ▶ just use the domain-specific representation for these

State-variable Representation 1

- state variables \mathbf{x} , \mathbf{y} for location, and \mathbf{u} , \mathbf{v} for velocity, values are integers
- action template:
 move(δ_u, δ_v) with $\text{Range}(\delta_u) = \text{Range}(\delta_v) = \{-1, 0, 1\}$
 pre: **safe**(δ_u, δ_v)
 - where **safe**(δ_u, δ_v) is a procedure to test whether the edge
 $((\mathbf{x}, \mathbf{y}), (\mathbf{x} + \mathbf{u} + \delta_u, \mathbf{y} + \mathbf{v} + \delta_v))$ intersects any of the walls
 eff: $\mathbf{x} \leftarrow \mathbf{x} + \mathbf{u} + \delta_u, \mathbf{y} \leftarrow \mathbf{y} + \mathbf{v} + \delta_v, \mathbf{u} \leftarrow \mathbf{u} + \delta_u, \mathbf{v} \leftarrow \mathbf{v} + \delta_v$
- Initial state, if (x_0, y_0) is the starting point:
 - ▶ $s_0 = \{\mathbf{x} = x_0, \mathbf{y} = y_0, \mathbf{u} = 0, \mathbf{v} = 0\}$
- Goal test, if the finish line is $((q, r), (q', r'))$:
 - ▶ $\min(q, q') \leq \mathbf{x} \leq \max(q, q'), \min(r, r') \leq \mathbf{y} \leq \max(r, r'), \mathbf{u} = \mathbf{v} = 0$
 - ▶ OK since we're allowing the goal test to be an arbitrary formula

State-variable Representation 2

- state variables \mathbf{p} , \mathbf{z} for location and velocity, each is a pair of integers

- action template:

move (δ_u, δ_v) with $\text{Range}(\delta_u) = \text{Range}(\delta_v) = \{-1, 0, 1\}$

pre: **safe** (δ_u, δ_v)

- where **safe** (δ_u, δ_v) is a procedure to test whether the edge $(\mathbf{p}, \mathbf{p} + \mathbf{z} + (\delta_u, \delta_v))$ intersects any of the walls

eff: $\mathbf{p} \leftarrow \mathbf{p} + \mathbf{z} + (\delta_u, \delta_v)$, $\mathbf{z} \leftarrow \mathbf{z} + (\delta_u, \delta_v)$

- Initial state, if (x_0, y_0) is the starting point:

► $s_0 = \{\mathbf{p} = (x_0, y_0), \mathbf{z} = (0, 0)\}$

- Goal test, if the finish line is $((q, r), (q', r'))$:

► $\min(q, q') \leq \mathbf{p}[0] \leq \max(q, q')$, $\min(r, r') \leq \mathbf{p}[1] \leq \max(r, r')$, $\mathbf{z} = (0, 0)$

State-variable Representation 3

- state variable \mathbf{s} for current state, a 4-tuple of integers

- action template:

move(δ_u, δ_v) with $\text{Range}(\delta_u) = \text{Range}(\delta_v) = \{-1, 0, 1\}$

pre: **safe**(δ_u, δ_v)

- where **safe**(δ_u, δ_v) is a procedure to test whether the edge
(($\mathbf{s}[0], \mathbf{s}[1]$), ($\mathbf{s}[0] + \mathbf{s}[2] + \delta_u, \mathbf{s}[1] + \mathbf{s}[3] + \delta_v$)) intersects any walls

eff: $\mathbf{s} \leftarrow (\mathbf{s}[0] + \mathbf{s}[2] + \delta_u, \mathbf{s}[1] + \mathbf{s}[3] + \delta_v, \mathbf{s}[2] + \delta_u, \mathbf{s}[3] + \delta_v)$

- Initial state, if (x_0, y_0) is the starting point:

► $s_0 = \{\mathbf{s} = (x_0, y_0, 0, 0)\}$

- Goal test, if the finish line is $((q, r), (q', r'))$:

$\min(q, q') \leq \mathbf{s}[0] \leq \max(q, q'), \min(r, r') \leq \mathbf{s}[1] \leq \max(r, r'), \mathbf{s}[2] = \mathbf{s}[3] = 0$

h^{FF} is better with some reps. than others

Suppose we start at $s_0 = ((1, 1), (0, 0))$, and choose $(u', v') = (1, 1)$

- Representation 1:
 - ▶ $\gamma^+(s_0, a) = \{\mathbf{x} = 1, \mathbf{x} = 2, \mathbf{y} = 1, \mathbf{y} = 2, \mathbf{u} = 0, \mathbf{u} = 1, \mathbf{v} = 0, \mathbf{v} = 1\}$
 - 16 possible states
 - ▶ RPG will converge quickly but probably won't be very informative
- Representation 2:
 - ▶ $\gamma^+(s_0, a) = \{\mathbf{p} = (1, 1), \mathbf{p} = (2, 2), \mathbf{z} = (0, 0), \mathbf{z} = (1, 1)\}$
 - 4 possible states; I think this representation will work OK
- Representation 3:
 - ▶ $s_0 = \{\mathbf{s} = (1, 1, 0, 0)\}$
 - ▶ $\gamma^+(s_0, a) = \{\mathbf{s} = (1, 1, 0, 0), \mathbf{s} = (2, 2, 1, 1)\}$
 - 2 possible states
 - ▶ RPG will do breadth-first search of the entire planning problem

Things I'll provide

I'll post a zip archive that includes the following:

- Domain-independent code (in Python 3.6):
 - ▶ `fsearch.py` – forward search algorithm
 - can do DFS, BFS, uniform-cost search, A*, and GBFS
 - has hooks for calling user-supplied code to draw search spaces
- Domain-specific code
 - ▶ `tdraw.py` – code to draw search spaces for racetrack problems
 - ▶ `racetrack.py` – code to run `fsearch.py` on racetrack problems
 - ▶ `maketrack.py` – Code to generate random racetrack problems
 - ▶ `sample_probs.py` – Some racetrack problems I generated by hand
 - ▶ `heuristics.py` – Some domain-specific heuristic functions
- `run_tests.bash` – a customizable **bash** script for running experiments

Here are some details ...

Contents of `fsearch.py`

Domain-independent forward-search algorithm

- `main(s0, next_states, goal_test, strategy, h=None, verbose=2, draw_edges=None)`
 - ▶ `s0` – initial state, in whatever representation you're using, e.g.,
 - for your FF heuristic, one of the state-variable representations
 - for the heuristic I give you, domain-specific representation
 - ▶ `next_states(s)` – function that returns the possible next states after `s`
 - ▶ `goal_test(s)` – function that returns `True` if `s` is a goal state, else `False`
 - ▶ `strategy` – one of `'bf'`, `'df'`, `'uc'`, `'gbf'`, `'a*'`
 - ▶ `h(s)` – heuristic function, should return an estimate of $h^*(s)$
 - ▶ `verbose` – one of 0, 1, 2, 3, 4
 - amount of verbosity in the output (see documentation in the file)
 - ▶ `draw_edges` – function to draw edges in the search space

Contents of `racetrack.py`

Code to run `fsearch.main` on `racetrack` problems

- `main(problem, strategy, h, verbose=0, draw=0, title='')`
 - ▶ `problem` – `[s0, finish_line, walls]`
 - ▶ `strategy` – one of `'bf'`, `'df'`, `'uc'`, `'gbf'`, `'a*'`
 - ▶ `h(s, f, w)` – heuristic function for `racetrack` problems
 - s = state, f = finish line, w = list of walls
 - `racetrack.py` converts this to the $h(s)$ function that `fsearch.main` needs (look at the code for this; you'll need something similar for `next_state`)
 - ▶ `verbose` – one of 0, 1, 2, 3, 4 (same as for `fsearch.py`)
 - ▶ `draw` – either 0 (draw nothing)
or 1 (draw problems, node expansions, solutions)
 - ▶ `title` – a title to use at the top of the graphics window
 - default is the names of the strategy and heuristic
- Some of the subroutines in `racetrack.py` may be useful ...

Contents of racetrack.py (continued)

- `intersect(e1,e2)` returns `True` if edges `e1` and `e2` intersect, `False` otherwise
 - `intersect([(0,0),(1,1)], [(0,1),(1,0)])` returns `True`
 - `intersect([(0,0),(0,1)], [(1,0),(1,1)])` returns `False`
 - `intersect([(0,0),(2,0)], [(0,0),(0,5)])` returns `True`
 - `intersect([(1,1),(6,6)], [(5,5),(8,8)])` returns `True`
 - `intersect([(1,1),(5,5)], [(6,6),(8,8)])` returns `False`

Basic idea (except for some special cases)


- ▶ Suppose $e1 = (p_1, p'_1)$, $e2 = (p_2, p'_2)$
 - ▶ Calculate the lines that contain the edges
 - $y = m_1x + b_1$; $y = m_2x + b_2$
 - ▶ If $m_1 = m_2$ and $b_1 \neq b_2$ then parallel, don't intersect
 - ▶ If $m_1 = m_2$ and $b_1 = b_2$ then collinear \Rightarrow check for overlap
 - ▶ If $m_1 \neq m_2$ then is intersection point contained in both edges?
- Can use this to implement the `goal_test` function for `fsearch.main`

Contents of racetrack.py (continued)

- `crash(e,walls)`
 - `e` is an edge
 - `walls` is a list of walls
 - True if `e` intersects a wall in `walls`, else False

- Example:

```
crash([(8,12),(10,13)],walls) returns False
crash([(8,12),(9,14)],walls) returns True
```



Contents of racetrack.py (continued)

- `children(state, walls)`


- ▶ state, list of walls
- ▶ Returns a list $[s_1, s_2, \dots, s_n]$
 - each s_i is a state that we can go to from `state` without crashing

- Example:

- ▶ current state is $((8, 12), (2, 2))$
- ▶ 9 possible states, 5 of them crash into the obstacle
- ▶ `children(((8,12),(2,2)), walls)` returns
 $[((9,13),(1,1)), ((10,13),(2,1)), ((11,13),(3,1)), ((11,14),(3,2))]$

- Can use this to implement the `next_states` function for `fsearch.main`

- ▶ That's a kluge since it's domain-specific, but I'll allow it anyway



Contents of `heuristics.py`

Three heuristic functions for the Racetrack domain:

- ▶ `h_edist(s, f, walls)` returns the Euclidean distance from *s* to the goal
 - can go in the wrong direction because it ignores walls
 - can overshoot because it ignores the number of moves needed to stop
- ▶ `h_esdist(s, f, walls)` is a modified version of `h_edist`
 - includes an estimate of howmany moves it will take to stop
- ▶ `h_walldist(s, f, walls):`

The first time it's called, for each gridpoint that's not inside a wall it will cache a rough estimate of the length of the shortest path to the finish line. The computation is done by a breadth-first search going backwards from the finish line, one gridpoint at a time.

On all subsequent calls, it will retrieve the cached value and add an estimate of how many moves will be needed to stop.

What you need to do

1. Write a Python function `ff1(s, f, walls)` that returns a minimal relaxed solution $\hat{\pi}$ for the racetrack problem $(s, f, walls)$.

Above, $\hat{\pi}$ is the same list $[\hat{a}_1, \hat{a}_2, \dots, \hat{a}_k]$ that *would be computed* by HFF (Algorithm 2.3) in the book or the pseudocode on page 54 of my `chap2b.pdf` lecture slides, if $(s, f, walls)$ were represented in state variable representation 1. However, you aren't required to use that pseudocode, and you aren't required to use state-variable representation (at least, not explicitly). It's OK for you to write your own *ad hoc* algorithm, as long as it returns correct answers.

In the minimal relaxed solution $\hat{\pi}$ that your algorithm returns, each \hat{a}_i should be a list of actions $\hat{a}_i = [a_{i1}, a_{i2}, \dots]$ in which each action a_{ij} is represented by just giving its arguments, e.g., use the pair $(1, -1)$ to represent the action `move(1, -1)`. Since all the actions are `move` actions, this representation is unambiguous.

2. Write a Python function `h1_ff1(s, g, walls)` that calls `ff1(s, g, walls)` to get the minimal relaxed solution $\hat{\pi}$, and returns the total number of actions in $\hat{\pi}$.

What you need to do (continued)

3. Write a Python function `ff2(s, f, walls)` that has the same description as `ff1(s, f, walls)`, with the phrase “state variable representation 1” replaced by “state variable representation 2”.
4. Write a Python function `h1_ff2(s, g, walls)` that calls `ff2(s, g, walls)` to get the minimal relaxed solution $\hat{\pi}$, and returns the total number of actions in $\hat{\pi}$.

Note: what “minimal” means

In the pseudocode for computing a minimal relaxed solution, if \hat{a}_i is a minimal set of actions that r-achieves \hat{g}_i , this means there’s no proper subset of \hat{a}_i that r-achieves \hat{g}_i . It’s OK if there’s a smaller set of actions \hat{a}'_i that r-achieves \hat{g}_i , as long as \hat{a}'_i isn’t a subset of \hat{a}_i .

There may be several minimal sets of actions that r-achieve \hat{g}_i , and they may have different sizes. Just take the first one you find; it doesn’t matter whether or not it’s the smallest one.

What you need to do (continued)

5. Do experiments to measure GBFS's performance as a function of problem size, using each of the following heuristics:
 - `h_ff1`
 - `h_ff2`
 - `h_esdist`
 - `h_walldist`
 - Later we'll give some guidelines about what ranges of problem size to use
5. Write a report giving the results of your experiments
 - Include the plots and table described on the next two pages
 - For each plot or table, tell what you can conclude from it and why
 - Format:
 - US letter paper, single column, 1-inch margins on all sides
 - Font size at least 11pt
 - At most 3 pages (or 4 if you do the extra-credit part)

Things to include in the report

- (a) a semi-log plot that shows, for each heuristic, the total CPU time for GBFS as a function of problem size.
 - ▶ Each data point: average of ≥ 10 randomly generated problems
- (b) Same as (a), but instead of CPU time, show number of nodes generated
- (c) A scatter plot for a subset of the Racetrack problems in `sample_probs.py` (I'll post a modified version of the file)
 - ▶ For each problem, plot a point (x, y) , where
 - x = number of nodes GBFS generates using `h_walldist`
 - y = number of nodes GBFS generates using `h_ff2`
- (e) A table showing the exact numbers in the scatter plot:
 - Column 1: problem name
 - Column 2: number of nodes GBFS generates using `h_walldist`
 - Column 3: number of nodes GBFS generates using `h_ff2`
 - ▶ The problems should be in the same order as in the `sample_probs.py` file

Grading

- Evaluation criteria:
 - ▶ 35% correctness: – whether your heuristic works correctly, whether your submission follows the instructions
 - ▶ 15% programming style – see the following
 - Style guide: <https://www.python.org/dev/peps/pep-0008/>
 - Python essays: <https://www.python.org/doc/essays/>
 - ▶ 15% documentation
 - Docstrings at start of file and in each function; comments elsewhere
 - ▶ 35% on the report itself
 - Adequacy of your experiments, statistical significance, clarity of presentation, quality of conclusions
- Extra credit:
 - ▶ Develop a better (for `h_ff`) state-variable representation for the Racetrack domain, and include it in your experiments
 - ▶ Caveat: I'm not sure whether that's feasible